

Best correct GPU kernel speedup per model (x = never produced correct output)

